

ML Project 5 - Health Care

February 16, 2023

1 ML Project 5 - Health Care

1.1 Task to be performed:

1. Preliminary analysis:
 - a. Perform preliminary data inspection and report the findings on the structure of the data, missing values, duplicates, etc.
 - b. Based on these findings, remove duplicates (if any) and treat missing values using an appropriate strategy
2. Prepare a report about the data explaining the distribution of the disease and the related factors using the steps listed below:
 - a. Get a preliminary statistical summary of the data and explore the measures of central tendencies and spread of the data
 - b. Identify the data variables which are categorical and describe and explore these variables using the appropriate tools, such as count plot
 - c. Study the occurrence of CVD across the Age category
 - d. Study the composition of all patients with respect to the Sex category
 - e. Study if one can detect heart attacks based on anomalies in the resting blood pressure (trestbps) of a patient
 - f. Describe the relationship between cholesterol levels and a target variable
 - g. State what relationship exists between peak exercising and the occurrence of a heart attack
 - h. Check if thalassemia is a major cause of CVD
 - i. List how the other factors determine the occurrence of CVD
 - j. Use a pair plot to understand the relationship between all the given variables
3. Build a baseline model to predict the risk of a heart attack using a logistic regression and random forest and explore the results while using correlation analysis and logistic regression (leveraging standard error and p-values from statsmodels) for feature selection

```
[1]: import pandas as pd
import numpy as np
```

```
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

2 1. Preliminary analysis:

- a. Perform preliminary data inspection and report the findings on the structure of the data, missing values, duplicates, etc.
- b. Based on these findings, remove duplicates (if any) and treat missing values using an appropriate strategy.

```
[2]: df = pd.read_excel('1645792390_cep1_dataset.xlsx')
```

```
[3]: df
```

```
[3]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	\
0	63	1	3	145	233	1	0	150	0	2.3	
1	37	1	2	130	250	0	1	187	0	3.5	
2	41	0	1	130	204	0	0	172	0	1.4	
3	56	1	1	120	236	0	1	178	0	0.8	
4	57	0	0	120	354	0	1	163	1	0.6	
..			
298	57	0	0	140	241	0	1	123	1	0.2	
299	45	1	3	110	264	0	1	132	0	1.2	
300	68	1	0	144	193	1	1	141	0	3.4	
301	57	1	0	130	131	0	1	115	1	1.2	
302	57	0	1	130	236	0	0	174	0	0.0	

	slope	ca	thal	target
0	0	0	1	1
1	0	0	2	1
2	2	0	2	1
3	2	0	2	1
4	2	0	2	1
..
298	1	0	3	0
299	1	0	3	0
300	1	2	3	0
301	1	1	3	0
302	1	1	2	0

[303 rows x 14 columns]

```
[4]: df.shape
```

```
[4]: (303, 14)
```

```
[5]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         303 non-null    int64
1   sex         303 non-null    int64
2   cp          303 non-null    int64
3   trestbps    303 non-null    int64
4   chol        303 non-null    int64
5   fbs         303 non-null    int64
6   restecg     303 non-null    int64
7   thalach     303 non-null    int64
8   exang       303 non-null    int64
9   oldpeak     303 non-null    float64
10  slope       303 non-null    int64
11  ca          303 non-null    int64
12  thal        303 non-null    int64
13  target      303 non-null    int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB
```

```
[6]: df.describe()
```

```
[6]:
```

	age	sex	cp	trestbps	chol	fbs	\
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	
mean	54.366337	0.683168	0.966997	131.623762	246.264026	0.148515	
std	9.082101	0.466011	1.032052	17.538143	51.830751	0.356198	
min	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	
25%	47.500000	0.000000	0.000000	120.000000	211.000000	0.000000	
50%	55.000000	1.000000	1.000000	130.000000	240.000000	0.000000	
75%	61.000000	1.000000	2.000000	140.000000	274.500000	0.000000	
max	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000	

	restecg	thalach	exang	oldpeak	slope	ca	\
count	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	
mean	0.528053	149.646865	0.326733	1.039604	1.399340	0.729373	
std	0.525860	22.905161	0.469794	1.161075	0.616226	1.022606	
min	0.000000	71.000000	0.000000	0.000000	0.000000	0.000000	
25%	0.000000	133.500000	0.000000	0.000000	1.000000	0.000000	
50%	1.000000	153.000000	0.000000	0.800000	1.000000	0.000000	
75%	1.000000	166.000000	1.000000	1.600000	2.000000	1.000000	
max	2.000000	202.000000	1.000000	6.200000	2.000000	4.000000	

	thal	target
count	303.000000	303.000000

mean	2.313531	0.544554
std	0.612277	0.498835
min	0.000000	0.000000
25%	2.000000	0.000000
50%	2.000000	1.000000
75%	3.000000	1.000000
max	3.000000	1.000000

```
[7]: df.isnull().sum()
```

```
[7]: age          0
     sex          0
     cp           0
     trestbps     0
     chol         0
     fbs          0
     restecg      0
     thalach      0
     exang        0
     oldpeak      0
     slope        0
     ca           0
     thal         0
     target       0
     dtype: int64
```

```
[8]: df.dtypes
```

```
[8]: age          int64
     sex          int64
     cp           int64
     trestbps     int64
     chol         int64
     fbs          int64
     restecg      int64
     thalach      int64
     exang        int64
     oldpeak      float64
     slope        int64
     ca           int64
     thal         int64
     target       int64
     dtype: object
```

```
[9]: duplicate = df[df.duplicated()]
```

```
[10]: duplicate
```

```
[10]:      age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  \
164    38    1   2        138   175    0         1       173     0       0.0

      slope  ca  thal  target
164      2   4     2        1
```

```
[11]: df = df.drop_duplicates()
```

```
[12]: df.duplicated().any()
```

```
[12]: False
```

3 2. Prepare a report about the data explaining the distribution of the disease and the related factors using the steps listed below:

- a. Get a preliminary statistical summary of the data and explore the measures of central tendencies and spread of the data
- b. Identify the data variables which are categorical and describe and explore these variables using the appropriate tools, such as count plot
- c. Study the occurrence of CVD across the Age category
- d. Study the composition of all patients with respect to the Sex category
- e. Study if one can detect heart attacks based on anomalies in the resting blood pressure (trestbps) of a patient
- f. Describe the relationship between cholesterol levels and a target variable
- g. State what relationship exists between peak exercising and the occurrence of a heart attack
- h. Check if thalassemia is a major cause of CVD
- i. List how the other factors determine the occurrence of CVD
- j. Use a pair plot to understand the relationship between all the given variables

```
[13]: df.describe()
```

```
[13]:      age      sex      cp      trestbps      chol      fbs  \
count  302.00000  302.000000  302.000000  302.000000  302.000000  302.000000
mean    54.42053    0.682119    0.963576   131.602649   246.500000    0.149007
std      9.04797    0.466426    1.032044    17.563394    51.753489    0.356686
min     29.00000    0.000000    0.000000    94.000000   126.000000    0.000000
25%     48.00000    0.000000    0.000000   120.000000   211.000000    0.000000
50%     55.50000    1.000000    1.000000   130.000000   240.500000    0.000000
75%     61.00000    1.000000    2.000000   140.000000   274.750000    0.000000
max     77.00000    1.000000    3.000000   200.000000   564.000000    1.000000
```

	restecg	thalach	exang	oldpeak	slope	ca \
count	302.000000	302.000000	302.000000	302.000000	302.000000	302.000000
mean	0.526490	149.569536	0.327815	1.043046	1.397351	0.718543
std	0.526027	22.903527	0.470196	1.161452	0.616274	1.006748
min	0.000000	71.000000	0.000000	0.000000	0.000000	0.000000
25%	0.000000	133.250000	0.000000	0.000000	1.000000	0.000000
50%	1.000000	152.500000	0.000000	0.800000	1.000000	0.000000
75%	1.000000	166.000000	1.000000	1.600000	2.000000	1.000000
max	2.000000	202.000000	1.000000	6.200000	2.000000	4.000000

	thal	target
count	302.000000	302.000000
mean	2.314570	0.543046
std	0.613026	0.498970
min	0.000000	0.000000
25%	2.000000	0.000000
50%	2.000000	1.000000
75%	3.000000	1.000000
max	3.000000	1.000000

```
[14]: df.mean(axis=0)
```

```
[14]: age          54.420530
sex           0.682119
cp            0.963576
trestbps     131.602649
chol         246.500000
fbs           0.149007
restecg       0.526490
thalach      149.569536
exang         0.327815
oldpeak       1.043046
slope         1.397351
ca            0.718543
thal          2.314570
target        0.543046
dtype: float64
```

```
[15]: df.median(axis=0)
```

```
[15]: age          55.5
sex           1.0
cp            1.0
trestbps     130.0
chol         240.5
fbs           0.0
restecg       1.0
```

```

thalach      152.5
exang        0.0
oldpeak      0.8
slope        1.0
ca           0.0
thal         2.0
target       1.0
dtype: float64

```

```
[16]: df.mode(axis=0).head(1)
```

```

[16]:      age  sex  cp  trestbps  chol  fbs  restecg  thalach  exang  oldpeak  \
0  58.0  1.0  0.0    120.0   197  0.0      1.0    162.0    0.0    0.0

      slope  ca  thal  target
0      2.0  0.0   2.0     1.0

```

```
[17]: cat_cols = ['sex', 'cp', 'fbs', 'restecg', 'exang', 'slope', 'ca', 'thal']
```

```
[18]: con_cols = ['age', 'trestbps', 'chol', 'thalach', 'oldpeak']
```

```
[19]: target_col = ['target']
```

```

[20]: print('Categorical Columns are :', cat_cols)
      print('Continuous Columns are :', con_cols)
      print('Target variable are :', target_col)

```

```

Categorical Columns are : ['sex', 'cp', 'fbs', 'restecg', 'exang', 'slope',
'ca', 'thal']
Continuous Columns are : ['age', 'trestbps', 'chol', 'thalach', 'oldpeak']
Target variable are : ['target']

```

```
[21]: df[con_cols].describe().transpose()
```

```

[21]:      count      mean      std    min    25%    50%    75%    max
age      302.0   54.420530   9.047970   29.0   48.00   55.5   61.00   77.0
trestbps 302.0  131.602649  17.563394   94.0  120.00  130.0  140.00  200.0
chol      302.0  246.500000  51.753489  126.0  211.00  240.5  274.75  564.0
thalach   302.0  149.569536  22.903527   71.0  133.25  152.5  166.00  202.0
oldpeak   302.0   1.043046   1.161452    0.0    0.00    0.8   1.60    6.2

```

```
[22]: df['age'].min()
```


```
[22]: 29
```

```
[23]: df['age'].max()
```

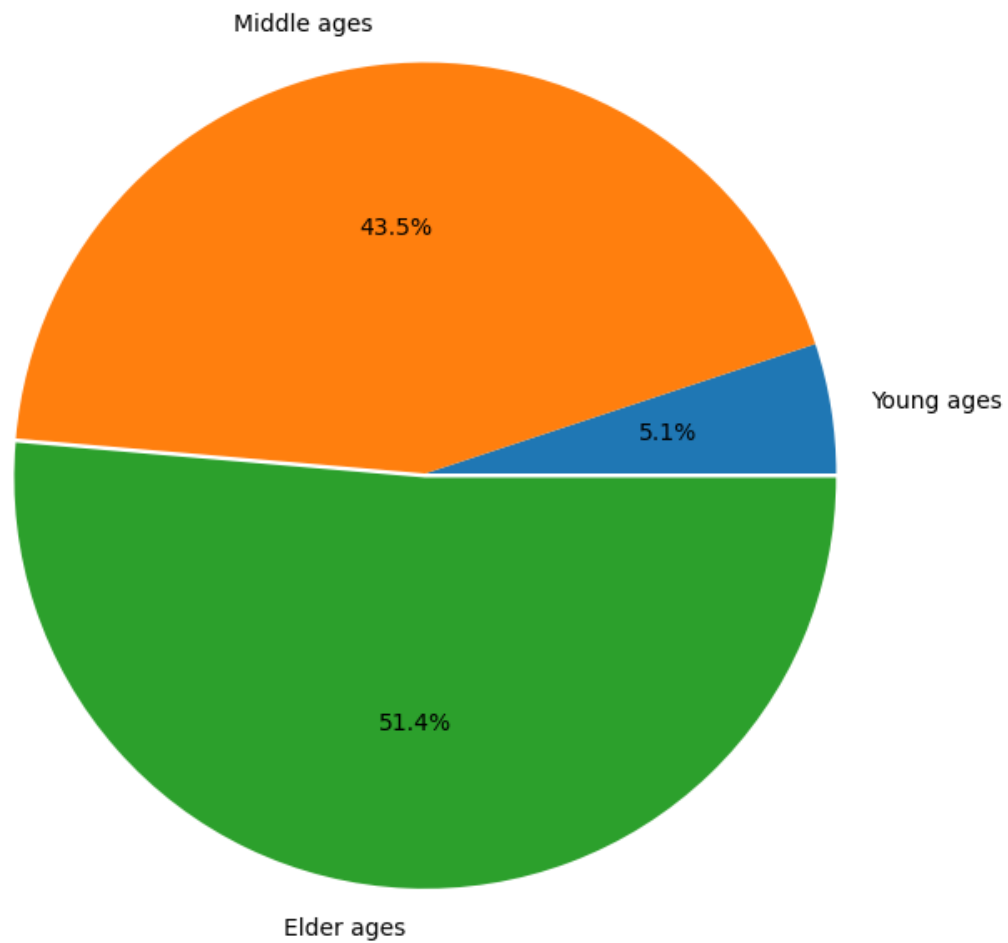
```
[23]: 77
```

```
[24]: Young = df[(df.age>=29) & (df.age<40)]
Middle = df[(df.age>=40) & (df.age<55)]
Elder = df[df.age>55]

plt.figure(figsize=(10,8))
sns.barplot(x=['Young ages', 'Middle ages', 'Elder_
↪ages'], y=[len(Young), len(Middle), len(Elder)])
plt.show()
```




```
[25]: plt.figure(figsize=(10,8))
plt.pie([len(Young), len(Middle), len(Elder)], labels=['Young ages', 'Middle_
↪ages', 'Elder ages'],
        explode=[0,0,0.01], autopct='%1.1f%%')
plt.show()
```

```
[26]: sns.countplot(df['sex'])  
plt.show()
```


```
C:\Users\Vinosh\anaconda3\lib\site-packages\seaborn\_decorators.py:36:  
FutureWarning: Pass the following variable as a keyword arg: x. From version  
0.12, the only valid positional argument will be `data`, and passing other  
arguments without an explicit keyword will result in an error or  
misinterpretation.  
warnings.warn(
```



```
[27]: sns.countplot(df['cp'])  
      plt.show()
```


```
C:\Users\Vinosh\anaconda3\lib\site-packages\seaborn\_decorators.py:36:  
FutureWarning: Pass the following variable as a keyword arg: x. From version  
0.12, the only valid positional argument will be `data`, and passing other  
arguments without an explicit keyword will result in an error or  
misinterpretation.  
    warnings.warn(  

```




```
[28]: sns.countplot(df['fbs'])  
      plt.show()
```

```
C:\Users\Vinosh\anaconda3\lib\site-packages\seaborn\_decorators.py:36:  
FutureWarning: Pass the following variable as a keyword arg: x. From version  
0.12, the only valid positional argument will be `data`, and passing other  
arguments without an explicit keyword will result in an error or  
misinterpretation.  
    warnings.warn(
```




```
[29]: sns.countplot(df['restecg'])  
plt.show()
```

```
C:\Users\Vinosh\anaconda3\lib\site-packages\seaborn\_decorators.py:36:  
FutureWarning: Pass the following variable as a keyword arg: x. From version  
0.12, the only valid positional argument will be `data`, and passing other  
arguments without an explicit keyword will result in an error or  
misinterpretation.  
warnings.warn(
```




```
[30]: sns.countplot(df['exang'])  
plt.show()
```

```
C:\Users\Vinosh\anaconda3\lib\site-packages\seaborn\_decorators.py:36:  
FutureWarning: Pass the following variable as a keyword arg: x. From version  
0.12, the only valid positional argument will be `data`, and passing other  
arguments without an explicit keyword will result in an error or  
misinterpretation.  
warnings.warn(
```



```
[31]: sns.countplot(df['slope'])  
      plt.show()
```


```
C:\Users\Vinosh\anaconda3\lib\site-packages\seaborn\_decorators.py:36:  
FutureWarning: Pass the following variable as a keyword arg: x. From version  
0.12, the only valid positional argument will be `data`, and passing other  
arguments without an explicit keyword will result in an error or  
misinterpretation.  
    warnings.warn(
```



```
[32]: sns.countplot(df['ca'])  
      plt.show()
```


```
C:\Users\Vinosh\anaconda3\lib\site-packages\seaborn\_decorators.py:36:  
FutureWarning: Pass the following variable as a keyword arg: x. From version  
0.12, the only valid positional argument will be `data`, and passing other  
arguments without an explicit keyword will result in an error or  
misinterpretation.  
    warnings.warn(  

```




```
[33]: sns.countplot(df['thal'])  
      plt.show()
```

```
C:\Users\Vinosh\anaconda3\lib\site-packages\seaborn\_decorators.py:36:  
FutureWarning: Pass the following variable as a keyword arg: x. From version  
0.12, the only valid positional argument will be `data`, and passing other  
arguments without an explicit keyword will result in an error or  
misinterpretation.  
    warnings.warn(
```





```
[34]: sns.boxenplot(df['age'])  
plt.show()  
  
sns.boxplot(df['age'])  
plt.show()
```

C:\Users\Vinosh\anaconda3\lib\site-packages\seaborn_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.
warnings.warn(




```
C:\Users\Vinosh\anaconda3\lib\site-packages\seaborn\_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.
  warnings.warn(
```




```
[35]: sns.boxenplot(df['trestbps'])  
plt.show()  
  
sns.boxplot(df['trestbps'])  
plt.show()
```

C:\Users\Vinosh\anaconda3\lib\site-packages\seaborn_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.
warnings.warn(




```
C:\Users\Vinosh\anaconda3\lib\site-packages\seaborn\_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.
  warnings.warn(
```




```
[36]: sns.boxenplot(df['chol'])  
plt.show()  
  
sns.boxplot(df['chol'])  
plt.show()
```

C:\Users\Vinosh\anaconda3\lib\site-packages\seaborn_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.
warnings.warn(




```
C:\Users\Vinosh\anaconda3\lib\site-packages\seaborn\_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.
  warnings.warn(
```




```
[37]: sns.boxenplot(df['thalach'])  
plt.show()  
  
sns.boxplot(df['thalach'])  
plt.show()
```

C:\Users\Vinosh\anaconda3\lib\site-packages\seaborn_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.
warnings.warn(




```
C:\Users\Vinosh\anaconda3\lib\site-packages\seaborn\_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.
  warnings.warn(
```





```
[38]: sns.boxenplot(df['oldpeak'])  
plt.show()  
  
sns.boxplot(df['oldpeak'])  
plt.show()
```

C:\Users\Vinosh\anaconda3\lib\site-packages\seaborn_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.
warnings.warn(




```
C:\Users\Vinosh\anaconda3\lib\site-packages\seaborn\_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.
  warnings.warn(
```



```
[39]: sns.countplot(df['target'])  
plt.show()
```

C:\Users\Vinosh\anaconda3\lib\site-packages\seaborn_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version
0.12, the only valid positional argument will be `data`, and passing other
arguments without an explicit keyword will result in an error or
misinterpretation.
warnings.warn(




```
[40]: df[con_cols].corr()
```


```
[40]:
```

	age	trestbps	chol	thalach	oldpeak
age	1.000000	0.283121	0.207216	-0.395235	0.206040
trestbps	0.283121	1.000000	0.125256	-0.048023	0.194600
chol	0.207216	0.125256	1.000000	-0.005308	0.050086
thalach	-0.395235	-0.048023	-0.005308	1.000000	-0.342201
oldpeak	0.206040	0.194600	0.050086	-0.342201	1.000000


```
[41]: plt.figure(figsize=(8,5))
sns.heatmap(df[con_cols].corr(),annot=True)
plt.show()
```




```
[42]: plt.figure(figsize=(10,5))
sns.kdeplot(df['age'],hue=df['target'],fill=True)
plt.show()
```




```
[43]: plt.figure(figsize=(10,5))
sns.kdeplot(df['trestbps'],hue=df['target'],fill=True)
plt.show()
```




```
[44]: plt.figure(figsize=(10,5))
sns.kdeplot(df['chol'],hue=df['target'],fill=True)
plt.show()
```




```
[45]: plt.figure(figsize=(10,5))
sns.kdeplot(df['thalach'],hue=df['target'],fill=True)
plt.show()
```




```
[46]: plt.figure(figsize=(10,5))
sns.kdeplot(df['oldpeak'],hue=df['target'],fill=True)
plt.show()
```




```
[47]: plt.figure(figsize=(10,5))
sns.kdeplot(df['cp'],hue=df['target'],fill=True)
plt.show()
```




```
[48]: plt.figure(figsize=(10,5))
sns.kdeplot(df['fbs'],hue=df['target'],fill=True)
plt.show()
```





```
[49]: plt.figure(figsize=(10,5))  
sns.kdeplot(df['restecg'],hue=df['target'],fill=True)  
plt.show()
```




```
[50]: plt.figure(figsize=(10,5))  
sns.kdeplot(df['exang'],hue=df['target'],fill=True)  
plt.show()
```



```
[51]: plt.figure(figsize=(10,5))
sns.kdeplot(df['slope'],hue=df['target'],fill=True)
plt.show()
```




```
[52]: plt.figure(figsize=(10,5))
sns.kdeplot(df['ca'],hue=df['target'],fill=True)
plt.show()
```




```
[53]: plt.figure(figsize=(10,5))
sns.countplot(df['sex'],hue=df['target'])
plt.show()
```


C:\Users\Vinosh\anaconda3\lib\site-packages\seaborn_decorators.py:36:
FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.
warnings.warn(




```
[54]: plt.figure(figsize=(10,5))
sns.kdeplot(df['thal'],hue=df['target'],fill=True)
plt.show()
```




```
[55]: sns.boxenplot(y=df['thalach'],x=df['target'])  
plt.show()  
  
sns.boxplot(y=df['thalach'],x=df['target'])  
plt.show()
```

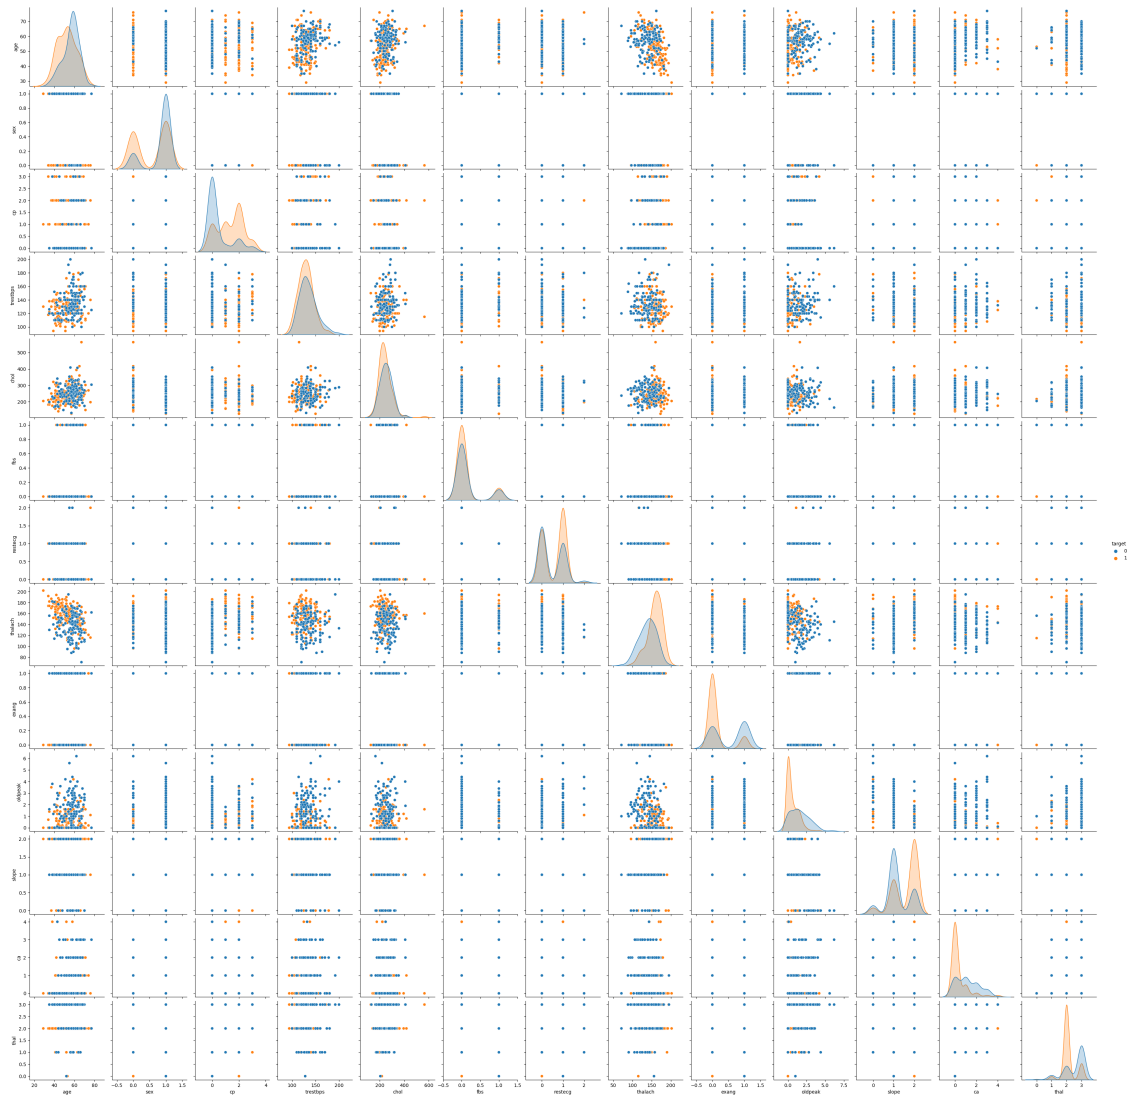




```
[56]: sns.stripplot(x=df['exang'],y=df['age'],hue=df['target'])  
plt.show()
```



```
[57]: sns.pairplot(df,hue='target')  
plt.show()
```



- 4 3. Build a baseline model to predict the risk of a heart attack using a logistic regression and random forest and explore the results while using correlation analysis and logistic regression (leveraging standard error and p-values from statsmodels) for feature selection

```
[58]: df1 = df
```

```
[59]: df1.head()
```



```
[59]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	\
0	63	1	3	145	233	1	0	150	0	2.3	0	
1	37	1	2	130	250	0	1	187	0	3.5	0	
2	41	0	1	130	204	0	0	172	0	1.4	2	
3	56	1	1	120	236	0	1	178	0	0.8	2	
4	57	0	0	120	354	0	1	163	1	0.6	2	

	ca	thal	target
0	0	1	1
1	0	2	1
2	0	2	1
3	0	2	1
4	0	2	1

```
[60]: cat_cols = ['sex','cp','fbs','restecg','exang','slope','ca','thal']
      con_cols = ['age','trestbps','chol','thalach','oldpeak']
```

```
[61]: df1 = pd.get_dummies(df1, columns=cat_cols)
```

```
[62]: df1.head()
```

```
[62]:
```

	age	trestbps	chol	thalach	oldpeak	target	sex_0	sex_1	cp_0	cp_1	\
0	63	145	233	150	2.3	1	0	1	0	0	
1	37	130	250	187	3.5	1	0	1	0	0	
2	41	130	204	172	1.4	1	1	0	0	1	
3	56	120	236	178	0.8	1	0	1	0	1	
4	57	120	354	163	0.6	1	1	0	1	0	

	...	slope_2	ca_0	ca_1	ca_2	ca_3	ca_4	thal_0	thal_1	thal_2	thal_3
0	...	0	1	0	0	0	0	0	1	0	0
1	...	0	1	0	0	0	0	0	0	1	0
2	...	1	1	0	0	0	0	0	0	1	0
3	...	1	1	0	0	0	0	0	0	1	0
4	...	1	1	0	0	0	0	0	0	1	0

[5 rows x 31 columns]

```
[63]: X = df1.drop(['target'],axis=1)
```

```
[64]: X
```

```
[64]:
```

	age	trestbps	chol	thalach	oldpeak	sex_0	sex_1	cp_0	cp_1	cp_2	\
0	63	145	233	150	2.3	0	1	0	0	0	
1	37	130	250	187	3.5	0	1	0	0	1	
2	41	130	204	172	1.4	1	0	0	1	0	
3	56	120	236	178	0.8	0	1	0	1	0	
4	57	120	354	163	0.6	1	0	1	0	0	
..	

298	57	140	241	123	0.2	1	0	1	0	0
299	45	110	264	132	1.2	0	1	0	0	0
300	68	144	193	141	3.4	0	1	1	0	0
301	57	130	131	115	1.2	0	1	1	0	0
302	57	130	236	174	0.0	1	0	0	1	0

	...	slope_2	ca_0	ca_1	ca_2	ca_3	ca_4	thal_0	thal_1	thal_2	\
0	...	0	1	0	0	0	0	0	1	0	
1	...	0	1	0	0	0	0	0	0	1	
2	...	1	1	0	0	0	0	0	0	1	
3	...	1	1	0	0	0	0	0	0	1	
4	...	1	1	0	0	0	0	0	0	1	
..			
298	...	0	1	0	0	0	0	0	0	0	
299	...	0	1	0	0	0	0	0	0	0	
300	...	0	0	0	1	0	0	0	0	0	
301	...	0	0	1	0	0	0	0	0	0	
302	...	0	0	1	0	0	0	0	0	1	

	thal_3
0	0
1	0
2	0
3	0
4	0
..	...
298	1
299	1
300	1
301	1
302	0

[302 rows x 30 columns]

```
[65]: y = df1['target']
```

```
[66]: y
```

```
[66]: 0      1
      1      1
      2      1
      3      1
      4      1
      ..
      298    0
      299    0
      300    0
```

```

301    0
302    0
Name: target, Length: 302, dtype: int64

```

```
[67]: from sklearn.preprocessing import StandardScaler
```

```
[68]: std = StandardScaler()
```

```
[69]: X[con_cols] = std.fit_transform(X[con_cols])
```

```
[70]: X.head()
```

```
[70]:
```

	age	trestbps	chol	thalach	oldpeak	sex_0	sex_1	cp_0	cp_1	\
0	0.949794	0.764066	-0.261285	0.018826	1.084022	0	1	0	0	
1	-1.928548	-0.091401	0.067741	1.636979	2.118926	0	1	0	0	
2	-1.485726	-0.091401	-0.822564	0.980971	0.307844	1	0	0	1	
3	0.174856	-0.661712	-0.203222	1.243374	-0.209608	0	1	0	1	
4	0.285561	-0.661712	2.080602	0.587366	-0.382092	1	0	1	0	

	cp_2	...	slope_2	ca_0	ca_1	ca_2	ca_3	ca_4	thal_0	thal_1	thal_2	\
0	0	...	0	1	0	0	0	0	0	1	0	
1	1	...	0	1	0	0	0	0	0	0	1	
2	0	...	1	1	0	0	0	0	0	0	1	
3	0	...	1	1	0	0	0	0	0	0	1	
4	0	...	1	1	0	0	0	0	0	0	1	

	thal_3
0	0
1	0
2	0
3	0
4	0


```
[5 rows x 30 columns]
```

```
[71]: from sklearn.model_selection import train_test_split
```

```
[72]: X_train,X_test,y_train,y_test = train_test_split(X,y,test_size=0.
↪20,random_state=42)
```

```
[73]: from sklearn.linear_model import LogisticRegression
```

```
[74]: lr = LogisticRegression()
```

```
[75]: lr.fit(X_train,y_train)
```

```
[75]: LogisticRegression()
```

```
[76]: from sklearn.feature_selection import f_regression
```

```
[77]: reg_score = pd.DataFrame(f_regression(X_train,y_train))
      reg_score
```

```
[77]:
```

	0	1	2	3	4	5	\	
0	10.979172	5.790709	0.315221	4.680671e+01	5.107591e+01	23.169146		
1	0.001064	0.016871	0.575020	6.520819e-11	1.070977e-11	0.000003		
	6	7	8	9	...	20	\	
0	23.169146	6.881119e+01	12.878390	23.933569	...	4.136295e+01		
1	0.000003	7.945327e-15	0.000403	0.000002	...	6.822119e-10		
	21	22	23	24	25	26	27	\
0	6.764737e+01	17.902679	17.146113	8.001474	0.172190	0.825156	3.211829	
1	1.257446e-14	0.000033	0.000048	0.005071	0.678545	0.364592	0.074373	
	28	29						
0	9.351766e+01	7.780579e+01						
1	6.997584e-19	2.429851e-16						

[2 rows x 30 columns]

```
[78]: print('F-score :',reg_score[0:])
```

```
F-score :
```

	0	1	2	3	4	5	\	
0	10.979172	5.790709	0.315221	4.680671e+01	5.107591e+01	23.169146		
1	0.001064	0.016871	0.575020	6.520819e-11	1.070977e-11	0.000003		
	6	7	8	9	...	20	\	
0	23.169146	6.881119e+01	12.878390	23.933569	...	4.136295e+01		
1	0.000003	7.945327e-15	0.000403	0.000002	...	6.822119e-10		
	21	22	23	24	25	26	27	\
0	6.764737e+01	17.902679	17.146113	8.001474	0.172190	0.825156	3.211829	
1	1.257446e-14	0.000033	0.000048	0.005071	0.678545	0.364592	0.074373	
	28	29						
0	9.351766e+01	7.780579e+01						
1	6.997584e-19	2.429851e-16						

[2 rows x 30 columns]

```
[79]: print('P-score :',reg_score[1:])
```

```
P-score :
```

	0	1	2	3	4	5	\
0	10.979172	5.790709	0.315221	4.680671e+01	5.107591e+01	23.169146	
1	0.001064	0.016871	0.575020	6.520819e-11	1.070977e-11	0.000003	

```

1  0.001064  0.016871  0.57502  6.520819e-11  1.070977e-11  0.000003

           6           7           8           9  ...           20  \
1  0.000003  7.945327e-15  0.000403  0.000002  ...  6.822119e-10

           21           22           23           24           25           26           27  \
1  1.257446e-14  0.000033  0.000048  0.005071  0.678545  0.364592  0.074373

           28           29
1  6.997584e-19  2.429851e-16

```

[1 rows x 30 columns]

```
[80]: y_pred = lr.predict(X_test)
```

```
[81]: from sklearn.metrics import accuracy_score, confusion_matrix, \
      ↪ classification_report
```

```
[82]: print('Accuracy Score :', accuracy_score(y_test, y_pred)*100)
```

Accuracy Score : 88.52459016393442

```
[83]: print(confusion_matrix(y_test, y_pred))
```

```

[[27  2]
 [ 5 27]]

```

```
[84]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.84	0.93	0.89	29
1	0.93	0.84	0.89	32
accuracy			0.89	61
macro avg	0.89	0.89	0.89	61
weighted avg	0.89	0.89	0.89	61

Accuracy score from Logistic Regression is 88.52%

```
[85]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.
      ↪ 20, random_state=42)
```

```
[86]: from sklearn.ensemble import RandomForestClassifier
```

```
[87]: rdc = RandomForestClassifier()
```

```
[88]: rdc.fit(X_train, y_train)
```

```
[88]: RandomForestClassifier()
```

```
[89]: y_pred1 = rdc.predict(X_test)
```

```
[90]: print('Accuracy Score :',accuracy_score(y_test,y_pred1)*100)
```

```
Accuracy Score : 85.24590163934425
```

```
[91]: print(confusion_matrix(y_test,y_pred1))
```

```
[[26  3]
 [ 6 26]]
```

```
[92]: print(classification_report(y_test,y_pred1))
```

	precision	recall	f1-score	support
0	0.81	0.90	0.85	29
1	0.90	0.81	0.85	32
accuracy			0.85	61
macro avg	0.85	0.85	0.85	61
weighted avg	0.86	0.85	0.85	61

Accuracy score from Random Forest is 85.23%

```
[ ]:
```